

gui_documentation

MIDAS GUI — User Documentation

Version: 1.0.0 **Application:** `midas-gui` (or `python -m midas_gui`) **Backends:**

`midas_calibrate_v2`, `midas_integrate_v2`, `midas_calibrate`, `midas_hkls`,

`midas_distortion`

Last updated: 2026-07-28 (Data Viewer: tilt-aware radial profile without a calibration file, ring 2 θ -range/thickness/live-button fixes, Save calibration as JSON/params/PONI, radial-plot X-axis floored at 0 and zoom preserved across parameter changes, Ctrl+S overwrites the loaded/saved session file with a new Ctrl+Shift+S "Save As")

Maintenance: keep this document in sync with the code — whenever the workflow or a tab's controls change, update the relevant section here in the same change.

Tab status: Tabs **0–4** (Data Viewer, Mask Builder, Calibrate, Calib. Refinement, Batch Integrate) are **verified** and ready to use. Tabs **5–8** (Corrections & Physics, PDF Analysis, Texture, Results & Export) are a **work in progress** and will be updated in the coming weeks.

Table of Contents

1. [Overview and Architecture](#)
2. [Getting Started](#)
3. [Data Viewer](#)
4. [Mask Builder](#)
5. [Calibrate](#)
6. [Calibration Refinement](#)
7. [Batch Integrate](#)
8. [Corrections & Physics](#)
9. [PDF Analysis](#)
10. [Texture / Pole Figure](#)
11. [Pump Probe \(time-resolved / TR-XRD\)](#)
12. [Results & Export](#)
13. [Common UI Conventions](#)
14. [Packaging, Deployment & Diagnostics](#)
15. [Configuration & Defaults](#)
16. [File ▶ Save/Load GUI State](#)

1. Overview and Architecture

The MIDAS GUI is a modular PyQt5 desktop application (up to 10 tabs) that exposes the scientific capability of the `midas_calibrate_v2` and `midas_integrate_v2` packages through a structured workflow. The intended order of use is:

```
Data Viewer (inspect) → Mask Builder → Calibrate → [Refine] → Batch Integrate
                                     → Corrections preview
                                     → PDF Analysis
                                     → Texture / Pole Figure
                                     → Pump Probe (TR-XRD)
                                     → Results & Export
```

Modular tabs. Data Viewer, Mask Builder, Calibrate and Batch Integrate are always shown; the remaining tabs (Calibration Refinement, Corrections, PDF Analysis, Texture, Pump Probe, Results & Export) are optional and can be shown/hidden from **Settings ▶ Preferences ▶ Tabs**. By default only **Calib. Refinement** and **Pump Probe** are shown; Corrections, PDF Analysis, Texture and Results & Export ship **hidden** — turn them on when you need them. The choice is saved per-user (`ui.visible_tabs`) and applies immediately — see §15. Hidden tabs are only removed from the tab bar; they stay constructed, so cross-tab wiring and state are preserved.

Cross-tab shared state. When Tab 2 (Calibrate) produces a result it is automatically propagated to all downstream tabs. When Tab 1 (Mask Builder) computes a mask it is sent to the consuming tabs. Tab 3 (Refinement), if applied, re-broadcasts the refined geometry. No manual copying is needed.

All heavy computation runs off the GUI thread. Every operation that touches a MIDAS package (calibration, integration, mask training, PDF transform, gain training, field averaging, drift fitting) runs in a background `QThread` worker; the GUI stays responsive and log lines stream in real time.

Package layout.

```
midas_gui/
├─ app.py           MainWindow, dark theme, crash diagnostics, main()
├─ __main__.py      enables `python -m midas_gui`
├─ style.py         Dioplas-inspired QSS theme + layout helpers
├─ constants.py     calibrants, colormaps, dtype sentinels, default paths
├─ helpers.py       image IO, geometry parsing, spec building, no-scroll widgets
├─ widgets.py       ImageViewer / PickableImageViewer / ProfileViewer / FieldSelecto
├─ workers.py       all QThread workers + the integration core
├─ calib.py         calibration-pipeline dispatch
```

```
|— dialogs.py      save dialogs
|— tab_*.py        one module per tab
```

Layout pattern. The four analysis tabs — **0 Data Viewer**, **2 Calibrate**, **3 Calib. Refinement**, **4 Batch Integrate** — use a **three-panel layout**:

```
[ Data Loader | Parameters | Display / results ]
```

- **Data Loader (left).** A shared `DataLoaderPanel` for selecting the five inputs — **Data**, **Dark**, **Bright**, **Background**, **Mask** — each as a single file, a folder, or an HDF5 dataset (a container dropdown appears for HDF5). Frame controls live here too (Tab 0 navigator, Tab 2/3 frame index, Tab 4 frame range + stride).
- **Parameters (middle).** The tab's analysis controls.
- **Display / results (right).** Image viewer, plots, and/or a log panel.

The three panels are separated by **draggable splitter handles** — drag the `Data Loader | Parameters` and `Parameters | Display` boundaries to rebalance widths to taste. Each panel has a minimum width; drag past it and the panel shows a scrollbar rather than clipping its controls.

The remaining tabs (1, 5–8) keep the classic two-panel layout.

Corrections & mask (all four analysis tabs). Dark is averaged then subtracted; Bright is averaged then flat-field **divided** or **subtracted** (your choice); Background is averaged then subtracted; every enabled Mask source (files/folders plus the auto-added "Tab 1 mask") is **unioned** into one composite mask that zeroes/ignores those pixels. Correction order: `(img - dark) → bright → - background → clip ≥ 0`.

2. Getting Started

`midas-gui` is **not on PyPI** — install it from source with conda.

```
git clone https://github.com/d-beniwal/MIDAS_GUI.git
cd MIDAS_GUI
conda env create -f environment.yml    # creates the 'midas-gui' env and installs
                                     the GUI
conda activate midas-gui
midas-gui                             # launch (equivalently: python -m midas_gui)
```

The GUI drives the MIDAS analysis backends (`midas-calibrate-v2`, `midas-integrate-v2`, `midas-calibrate`, `midas-hkls`, `midas-distortion`), which are also not on PyPI — point the `pip:` section of `environment.yml` at your MIDAS source before creating the environment (see the comments in that file).

Test data

Synthetic test data lives in the repo-root `test_data/` folder (git-ignored):

- `calibrant_ceria.tif` / `.h5` — CeO₂ calibrant (Eiger2 500K, 1028×512, 75 μm pixels)
- `nickel_tifs/` — 10-frame Ni scan, lattice expanding +0.1 %/frame
- `nickel_stack.h5` — the same scan as one HDF5 stack (`exchange/data`)

Default geometry: $\lambda = 0.39 \text{ \AA}$, Lsd = 121 mm, pixel = 75 μm, BC = (10, 10) px. Every tab's default file paths point at this data and **auto-load on startup when present**, so the app is usable immediately after checkout.

3. Tab 0 — Data Viewer

Purpose: inspect detector frames and do a quick radial integration — geometry-free (circle binning) by default, or a full tilt/distortion-aware integration when a calibration file is loaded. Produces no shared state for other tabs.

Data Loader panel (left)

Data, Dark, Bright, Background and Mask are all selected in the shared **Data Loader panel** (see §1): each accepts a file / folder / HDF5-dataset. The **Data** card here holds the **frame navigator** (slider, spin box, ◀/▶); **zoom/pan is preserved** as you step through frames (the view only auto-frames on a fresh load). Dark/bright/background and the composite mask are applied to the displayed image and to the radial integration. **Pan and zoom are bounded to the image** (roughly half an image-width of margin on each side) so scrolling/dragging cannot wander off into empty space or zoom out indefinitely; the bottom-left **A** (auto-range) button re-fits the image without leaving zoom "stuck" tracking the mouse.

Projection card

Collapse a stack to one image: **Max** (hot-pixel hunting), **Sum** (long-exposure equivalent), or **Average** (noise reduction), along a chosen axis (0 = across frames). **Skip frames** (default 1) ignores that many leading frames before projecting (e.g. 1 drops the first frame, 4 drops the first four) — useful when the opening frames are detector warm-up / shutter-transient exposures.

Calibration card (optional)

Load geometry from a **calibration** `.json`, a MIDAS `paramstest.txt`, or a pyFAI `.poni` (auto-detected). It fills BC, Lsd, pixel size and wavelength, unchecks "Beam centre = image centre", and refreshes the overlay and radial plot.

When a calibration file carries the **full geometry (tilts + distortion)**, the radial integration switches from simple concentric-circle binning to a **proper MIDAS-engine integration** that maps every pixel through the calibrated tilts and distortion (the same core as Batch Integrate) — so ring positions/intensities are geometry-correct, not just distance-from-beam-centre. The card's status line reports which mode is active. The binning geometry is built once and reused across frames (a fast `hard` kernel keeps the preview responsive); without a full-geometry file, the fast circle binning is used. If no calibration file is loaded but the Ring-simulation card's **ty/tz** tilt fields are non-zero, the radial integration still runs the tilt-aware MIDAS engine using a geometry built live from those fields (BC, Lsd, pixel size, wavelength) — so dialing in tilts here without a calibration file already produces a tilt-corrected profile, not just a tilt-shaped ring overlay.

Save JSON / Save params (.txt) / Save PONI (below the loader) export whatever geometry is currently in effect — the loaded calibration file if any, otherwise the one synthesized from the Ring-simulation widgets above — as a calibration file you can reload here, on the Calibrate tab, or in Batch Integrate. **Save params (.txt)** writes a standalone MIDAS `paramstest.txt` ; **Save PONI** writes a pyFAI `.poni` (note: PONI's Rot1–3 convention cannot represent MIDAS's tx/ty/tz tilts, so tilts are **not** included in a `.poni` export — use JSON or `.txt` to keep them). Clicking any of the three before an image is loaded warns instead of writing a file (there is no detector size to write).

Intensity range card (radial integration)

Field	Description
Exclude out-of-range pixels	When on, pixels $\leq$ min or $>$ max are drawn as a red overlay and excluded from the radial integration (removes gaps / hot / overflow).
pixel $\leq$ / pixel $>$	Lower / upper bounds. On load, the upper bound auto-fills to max(99.99th percentile, 100000) .

Ring simulation card

Overlays simulated Debye-Scherrer rings to check geometry.

- **Material** dropdown (CeO₂, LaB₆, Si, Al₂O₃, Cu, Ni, FCC- γ Fe, BCC- α Fe, Au, Ag, Pt, W, Ti) or **Custom**.
- Lattice entry is compact: **a**, **b**, **c** on one row and **α** , **β** , **γ** on the next, plus **SG #**. A **Cubic (a=b=c, $\alpha=\beta=\gamma=90^\circ$)** checkbox lets you enter only **a** for cubic crystals (b, c mirror a; angles fixed at 90°). Lattice fields are editable only for Custom.
- Geometry: λ , max 2 θ , Lsd, pixel size, and beam centre (auto = image centre, or manual BC_y/BC_z). The λ label is clickable (underlined) — click it for a menu of common K-edge foils (Pr, Sm, Yb, Lu, Hf, Ta, W, Re, Pt, Au, Pb, Bi) that sets λ to that element's K absorption-edge wavelength. The same clickable- λ menu is on the Calibrate and PDF tabs. The **px** label is likewise clickable — a menu of common detectors (GE 200 μ m, Varex

150 μm , Pilatus 172 μm , Eiger 75 μm) sets the pixel size (also on the Calibrate tab, where it sets both pxY and pxZ).

- Every field in this card steps by a fixed amount per up/down-arrow click (λ 0.01 $\text{\AA}$, max 2θ 1°, Lsd 1 mm, pixel 0.1 μm , BC_y/BC_z 1 px, ty/tz 0.1°) — customizable from **Settings ▶ Preferences ▶ Data Viewer**.
- **ty / tz** (detector tilt about the Y/Z axes, degrees) sit right below BC_y/BC_z — when non-zero, the simulated rings are forward-projected through the tilt geometry instead of drawn as plain circles, so the overlay shows the same non-circular ring shape a tilted detector actually produces. Leaving both at 0° reproduces the previous plain-circle rendering exactly. (Rotation about the beam axis, **tx**, leaves a full ring's shape unchanged, so it isn't exposed here.)
- **Ring thickness** spin box (0.5–10 px) sets the line width of the simulated-ring overlay on the image; redraws immediately as you change it.
- **Show rings / Labels** toggles; **Simulate rings** is a **live toggle** (not a one-shot click) — while on, the button turns **green** as a visual cue, and the overlay + hkl table recompute automatically whenever material, lattice constants, space group, or geometry (λ /Lsd/px/max 2θ) change. Beam-centre and ty/tz edits still just reposition the existing rings (their 2θ values don't depend on BC or tilt). Rings are drawn out to the full **max 2θ** you set, regardless of how far that places them from the beam centre (rings landing off-canvas simply aren't visible — they aren't silently dropped at some fixed pixel-radius cutoff).
- → **Send geometry to Calibrate** copies λ , pixel size, Lsd and beam centre into the Calibrate tab's detector + seed fields (the Calibrate tab has a matching **← Data Viewer** button that pulls the same values).

Live Data card (*experimental*)

Sits **above the Data card**, collapsed by default behind its own title-bar checkbox — check it to reveal the live-PV controls, uncheck to hide them again (unchecking also stops an active stream, so a hidden card can never be left silently connected). Subscribes directly to an EPICS PVA detector-image PV (NTNDArray) and renders each frame **inline in the Data Viewer's own image pane** — the same view used for files/folders/HDF5 stacks. Because it feeds the normal frame pipeline, all existing Data Viewer analysis keeps working live: dark/bright/background/mask corrections, the intensity-range mask, beam-centre picking, and the radial integration plot all update as new frames arrive.

- **Dependency:** `pvapy` (the EPICS pvAccess client library) is a required, pinned dependency (`pyproject.toml` / `environment.yml`), installed automatically with the rest of the GUI's stack — no separate extra to install. If somehow missing from the active environment, **Start** shows an install hint instead of failing outright.
- **Live PV** field is an editable dropdown: pick a known device by name (the list comes from **Preferences ▶ Devices**, see below) to fill in its full PV automatically, or type any other PV by hand (placeholder shows an example, `20IDFF:Pva1:Image`). **Start / Stop** buttons and a

status line (stopped / waiting for PV / connected / streaming with frame id / error). GUI updates are throttled to the tab's existing ~16 fps debounce, so a fast PV update rate doesn't overwhelm the interface.

- **Stopping live streaming does not auto-reload the previous file/folder/HDF5 source** — the Data card below gets a small  **Reload** button next to its path field for exactly this: click it after Stop to restore the static data that was loaded before streaming started.
- **Ring overlay is static per-frame while streaming:** rings drawn by "Simulate rings" (Ring simulation card) keep rendering on top of each new live frame at their already-computed geometry — arrival of a new frame does **not** by itself trigger a recompute (ring radii don't depend on frame data). With Simulate rings' live toggle on, changing a material/lattice/geometry field still recomputes and redraws immediately. **Stop** leaves the last received frame on screen. Closing midas-gui also stops any running stream.
- **Colormap/level changes persist across frames:** dragging the histogram's level range (or the cmap dropdown) is remembered and reapplied to every new incoming frame instead of being reset to the vmin%/vmax% percentile defaults; editing vmin%/vmax% explicitly switches back to auto-levels.

Beam-centre picking (on the image)

The image viewer has **Pick BC** (single click sets the beam centre) and **Pick Ring** (click ≥ 3 points on a ring; a circle fit estimates the beam centre). Either updates BC_y/BC_z and re-runs the overlay + radial integration. Pick Ring's picked points and fitted circle are drawn in **blue**, distinct from the **amber** Simulate-rings overlay so the two aren't confused when both are visible on the same image.

Top-N brightest pixels (image toolbar)

A **Top-N pixels** toggle button (with an **N** spin box) sits on the image toolbar. When on, the **N highest-intensity pixels** of the current frame are marked with a crosshair inside a translucent circle (cyan) centred on each pixel, and the Intensity statistics panel switches to show the statistics of just those N pixels ("Top N pixels"). It follows frame changes while active. Clicking the button again removes the markers and restores the normal statistics. Useful for quickly locating saturation / hot pixels.

Intensity statistics (left panel, bottom)

At the bottom of the Data-Loader panel, in a **draggable pane** below the loader cards — grab the splitter handle above the panel to make the statistics area taller or shorter. It holds a **histogram** of the intensity distribution (full range, log-y toggle) with a **textbox** beneath it reporting N (pixel count) and the **p70 / p90 / p99 / p99.9 / p99.99** percentiles — each with the **number of pixels above** that value. The histogram's lower-left corner is fixed at $x = 0$, $y = -2$ and both axes rescale to $(0, x_{\max}) / (-2, y_{\max})$ on every refresh (frame change, scope change, projection, new data). It reflects the **corrected** image (dark / bright / background) with masked pixels (file masks + the intensity-range mask) excluded, and updates live as any of those change. A scope

selector switches between the **current frame** (per the slider) and **All frames** (combined over the whole stack/folder). When a **Projection** is active the panel shows the projected image's statistics (the scope selector is disabled); when **Top-N pixels** is active it shows those pixels' statistics.

Radial integration plot (bottom-right)

Below the image is a live **azimuthal average about the beam centre** (`R bin` , `Integrate` , and an `Auto` toggle that recomputes on frame/BC/mask change) — geometry-free circle binning by default, or the tilt/distortion-aware MIDAS engine when a full geometry is in effect (a loaded calibration file, or non-zero ty/tz in the Ring-simulation card — see the Calibration card above). Peak/ring markers overlaid on the plot use the ring's true 2θ , so they line up with the profile in either mode. **Clicking a radius on the plot draws the matching ring (magenta) on the image.** Axis units switch between R (px) / 2θ / Q ; the **X-axis lower bound defaults to 0**.

The default view always auto-fits the current profile's X/Y extent — it never gets stuck showing a stale or unrelated range from an earlier profile. If you manually zoom or pan (drag, wheel-zoom, box-zoom, or the axis context menu), that exact view is **preserved across parameter changes** (new frame, changed BC/tilt/Lsd, etc.) instead of snapping back to full range every time the profile updates — the curve redraws in place under your current zoom. A manual zoom is only cleared when it would no longer make sense: switching the X-axis unit ($R/2\theta/Q$) drops the remembered X range, and toggling **Log Y** drops the remembered Y range (both because the old numbers belong to a different scale). **Pan and zoom are always bounded to the current profile's extent** (a margin around its X/Y range) so you cannot drag or scroll off into empty space. The splitter handle above this panel (between it and the image view) is wider than a default Qt splitter, to make it easier to grab.

4. Tab 1 — Mask Builder

Purpose: identify and exclude bad pixels. The final mask is a uint8 array (1 = bad) broadcast to Tabs 2, 3, 4, 5. Browsing an image or a mask loads it immediately. Pointing the **Image** field at an `.h5` file reveals a **Dataset** dropdown (auto-populated with the file's datasets and shapes, preferring a ≥ 2 -D dataset) — the same pattern used by the Stack field below it and by the Data Loader panels elsewhere in the app.

Section 1 · Threshold mask (always applied)

`pixel ≤ lower | pixel > upper` . The upper bound auto-fills from the data type on load (e.g. 1,048,575 for uint20 Eiger, 4,294,967,295 for uint32).

Section 2 · Statistical auto-mask

Two **independently selectable** methods — enable either or both:

- **Spatial outlier** — robust local-anomaly detection (5×5 median residual → 15×15 local MAD → Z-score → hot/dead/saturated gates). Uses a **temporal median** over the stack folder when provided (cleaner reference), else the single frame. Controls: `K_σ` (6.0), `Hot` (1.5), `Dead` (0.5).
- **Temporal constancy** — flags frozen pixels whose frame-to-frame std is below `Frozen × Q75(std)` (0.05); needs a stack of ≥2 frames.

A shared **stack source + stride** feeds both (the temporal median for spatial, the frame stack for temporal constancy). The stack can be a **folder / *.tif glob**, or a **single HDF5 file whose 3-D dataset is a time sequence of images** — for an `.h5` a **Dataset** selector appears (auto-populated, 3-D datasets preferred). *All σ / K_σ / n_σ fields in this tab accept any value — there is no upper limit.*

Section 3 · Spatial spike rejection

Laplacian high-pass single-pixel-spike detector; `n_σ` threshold (5.0).

Section 3b · Cosmic-ray rejection (temporal)

Per-pixel temporal σ -clip across a ≥3-frame stack; anomalies OR'd across frames. `n_σ` (5.0).

Section 4 · Calibration-based masks (need a Tab 2 result)

- **Azimuthal σ -clip** — flags (R, η) cells deviating from the azimuthal mean.
- **Learnable mask** — differentiable per-pixel mask trained against an η -uniformity loss (`steps`, `lr`, `sparsity`).

Draw mask tools

Interactive rectangle / oval / circle / polygon / annulus / point tools, live-draggable; **Apply shapes** → **mask** rasterises them (OR'd with the computed mask). **Clear shapes** removes them.

Save / Load

Save the combined mask as TIFF (0 = good, 1 = bad); loading a TIFF applies immediately.

5. Tab 2 — Calibrate

Purpose: determine detector geometry (Lsd, BC, tilts, distortion, wavelength) from a calibrant pattern.

Data Loader panel (left)

The calibrant frame (Data + Frame index), Dark, Bright, Background and Mask are selected in the shared **Data Loader panel** (see §1). Each Dark/Bright/Background field is averaged over a chosen index range; Bright offers **Flat-field divide** or **Subtract**; all mask sources (plus the auto-added Tab 1 mask) are unioned. Dark is passed to the calibration pipeline; bright/background are applied to the calibrant before calibration and post-calibration integration.

Detector, seed & Load calibration file

The **Detector & Calibrant** card sets λ , pixel size(s) and detector transforms; the **Initial seed** card sets the LM starting point (BC_y, BC_z, Lsd, and tilts tx/ty/tz). **Load calibration file...** (top of the Detector card) reads a MIDAS **paramtest .txt**, a calibration **.json**, or a pyFAI **.poni** (auto-detected) and fills λ , pixel size, and the seed BC + Lsd — a fast way to start from a previous calibration or a known geometry. (The synthetic test data ships

`test_data/calibration_synthetic.{json,txt,poni}` as a ready example.) ← **Data Viewer** (next to *Load calibration file...*) pulls λ , pixel size, Lsd, beam centre and the Data Viewer's ty/tz tilt fields straight from the Data Viewer tab into the same fields (BC and Lsd land in the seed card; ty/tz land in the seed tilt fields). A **Feed result back to seed** checkbox (on by default) copies the optimized BC / Lsd / tilts / distortion of each run back into the seed fields, so a follow-up run starts from the previous solution. *Seed tilts are honoured by the Four-stage / advanced pipelines; the One-shot / First-time paths seed tilts only if the installed backend exposes initial-tilt options.*

Average frames

For a multi-frame source (HDF5 / folder), **Average frames into a single image** builds the mean of a frame range and calibrates on that. **start / end (0 = all)** select the range and **skip** uses every Nth frame while averaging (e.g. skip = 2 averages every other frame). The card is disabled for single-frame sources; the preview updates live as the options change.

Pipeline selector

One-shot (default) · First-time · Four-stage · Bayesian (Laplace σ) · Joint-cake. *For trustworthy tilt/strain, prefer Four-stage or First-time — One-shot/Bayesian can report a spurious self-compensated tilt on weakly-tilted data.*

Refine flags

Which parameters vary: Lsd, BC, ty, tz, tx, Wavelength, plus a "Residual map" build toggle. The **Distortion (n/15)** checkbox has a companion ... button that opens a per-coefficient dialog: the 15 distortion coefficients are grouped by η -fold (isotropic radial + folds 1–6), and named **preset modes** (None · Isotropic only · Iso + up to 2-fold · Iso + up to 4-fold · All (15)) auto-select whole ladders. The checkbox label shows how many coefficients are selected. *Per-coefficient control applies to the Four-stage / advanced pipelines; the One-shot path refines all-or-none.* Advanced (E-M / LM iters, device, output dir) and Multi-panel detector groups are collapsible.

Live threshold slider

Zeroes calibration-image pixels below the slider value (background suppression for ring-finding); the preview updates instantly.

Pick BC / Pick Ring

Click the image to seed the beam centre (single click) or fit a ring (≥ 3 clicks); the Pick Ring points/fit are drawn in blue, distinct from the amber Simulate-rings overlay.

Predicted-ring overlay (image toolbar)

After a run, the calibrant's predicted ring radii are drawn in **lime** with a red/yellow beam-centre marker. **Show rings** toggles the overlay; **Corrected** redraws the same lime rings reshaped through the fitted tilt (tx/ty/tz) instead of as plain circles, so you can see how much the geometry actually deviates from an untilted detector — the rings still look like the same thin curves, just bent, rather than a separate scattered point-cloud.

Run / Abort

Run Calibration launches the worker; **Abort** terminates it and frees the slot so you can immediately start a new run (the calibration is one uninterruptible library call, so abort hard-terminates the worker thread rather than waiting).

Results (right panel, bottom tabs)

Radial Profile (with ring markers and an **Azim. avg** selector — see Tab 4), Ring Residuals bar chart, **Results**, and Log. Integration runs automatically after calibration. The **Results** tab shows the *complete* parameter set exactly as it is written to `paramstest.txt` (Lsd, BC, tx/ty/tz, the distortion coefficients, Parallax, Wavelength, px, NrPixelsY/Z, RhoD, SpaceGroup, LatticeConstant) as **plain text laid out in multiple columns** so the wide-but-short panel stays readable — no table widget. The distortion `p0-p14` slots are labelled with their coefficient names (iso_R2, a1, phi1, ...). A strain/timing line sits below. → **Send to Data Viewer** pushes the *full* calibrated geometry (λ , pixel size, Lsd, beam centre, **tilts and distortion**) into the Data Viewer tab — where it drives the tilt/distortion-aware radial integration, not just circle binning — the reverse of the Data Viewer's "→ Send geometry to Calibrate". The bottom tab area is fully resizable (drag the horizontal splitter) and the Log fills its tab.

Export

Save calibration.json and **Save paramstest.txt** (standalone or from a template).

6. Tab 3 — Calibration Refinement

Purpose: optional post-calibration geometry refinement against an **η -uniformity** criterion (rings should be azimuthally uniform). Uses a **derivative-free Nelder-Mead** optimiser (the differentiable integrator returns NaN geometry gradients at the beam-centre singularity on this build).

The sample frame and Dark/Bright/Background/Mask are selected in the shared **Data Loader panel** (see §1) — the refinement runs on the corrected image. Middle-panel controls: calibration source (Tab 2 result, or start-from/continue radios), parameters to refine (BC_y, BC_z, Lsd, ty, tz), optimiser + iterations. **Run Refinement** shows a live loss curve; **Apply** broadcasts the refined geometry downstream. If Apply is never clicked, the Tab 2 result flows through unchanged.

7. Tab 4 — Batch Integrate

Purpose: integrate a stack of frames into 1-D profiles using the calibrated geometry and optional corrections.

Data Loader panel (left)

The streaming **Data** source (folder/glob or HDF5 dataset) with **frame range + stride**, plus Dark/Bright/Background and Mask, are selected in the shared **Data Loader panel** (see §1). Dark/bright/background are applied per frame; all mask sources are unioned.

Calibration source (middle)

- **From Tab 2** — the calibration (or refined) result.
- **From file** — a calibration `.json`, MIDAS `paramstest.txt`, or pyFAI `.poni` (auto-detected; both GUI and MIDAS-pipeline json key styles supported). Entering a path auto-selects this option.

Calibration values (middle, read-only)

A **Calibration values** card shows the geometry actually in use — λ , Lsd (mm), BC_y/BC_z (px), tilts tx/ty/tz (°), pixel sizes, detector size, and a distortion summary (number of non-zero coefficients). It repopulates whenever a calibration arrives from Tab 2 (or Tab 3 refinement, or the Data Viewer via Tab 2) **and** when a calibration file path is entered, parsing the file directly so you can confirm the numbers before running. A note line reports the active source (or any file-read error).

Integration

Field	Description
Kernel	Hard (fastest) · Subpixel K=2 · Subpixel K=4 · Polygon (exact).

Field	Description
R bin (px) / η bin (°)	Radial and azimuthal bin sizes.
Azim. avg	How the (η, R) cake becomes a 1-D profile: Pixel-weighted (default) $\Sigma(\text{mean} \cdot \text{count}) / \Sigma(\text{count})$ — independent of η-bin size and robust to partial azimuthal coverage / off-detector beam centres ; or η-bin mean (legacy) — the unweighted mean of per-η-bin means, which can distort with a coarse η bin when the beam centre is off the detector.
Per-bin variance (σ)	Error model poisson / azimuthal / hybrid (ignored when corrections are on $\rightarrow \sigma = \sqrt{I}$).
Q-uniform bins	Integrate in R then rebin onto a uniform-Q grid (Qmin, Qmax, ΔQ).

Physics corrections

Polarization and solid-angle (pixel-domain, via `integrate_with_corrections`).

Monitor normalisation

Divide each processed frame's profile/σ by a per-frame scalar from a text file (one value per *processed* frame).

Drift correction (long scans)

Per-frame geometry interpolated from an anchor JSON (`{frame_idx: {Lsd, BC_y, BC_z}}`) with spline/linear/constant parametrization; **Fit trajectory** builds it. Each frame is then integrated with its own geometry.

Run / Abort / Clear

Start Integration / Abort. Abort first asks the worker to stop cleanly between frames (keeping frames already written); if it does not stop promptly it force-terminates and frees the slot. Completed frames' output files are preserved. **Clear results** removes the profiles/plots computed **this session** for the current data (waterfall, stacked profiles, and the integrated-frame tracking) so a fresh integration can start — it stops any active monitor but does **not** delete raw data or any files already written to disk.

Live folder monitoring (MONITOR)

The **MONITOR** button at the bottom of the left Data Loader panel (folder/glob sources only) starts a live watch: while active it turns **green** and the GUI polls the data folder for **new** TIFF

frames, integrating each one **as it appears** and adding it to the display. It does **not** re-run the whole batch — it reuses the already-built **detector map** (the binning geometry / pixel-count cakes; reused from a prior *Start Integration* run when the calibration, kernel, bins, mask and folder are unchanged, or built once on first use) and integrates **only the new files** (tracked by frame id, so frames already shown are skipped). New frames honour the current kernel, corrections, Dark/Bright/Background, mask and Q-uniform settings, and are saved to the output folder when a 1-D format is selected. Click MONITOR again to stop; starting a fresh *Start Integration* also stops it. (Distinct from **Monitor normalisation** above, which divides profiles by a scalar file.)

Output formats

CSV (R,I, σ) · XYE (2 θ) · FXYE (centideg) · DAT (Q) · HDF5 (full stack) · 2D-CSV ($\eta \times R$ cake). Right panel: live **Waterfall** and **Stacked profiles** — both have an x selector to show the axis in **R (px) / 2 θ (°) / Q ($\text{\AA}^{-1}$)** (converted from the run's calibration).

The **Stacked profiles** view is a publication-quality plot: each curve tags its **source file name inline, just below the curve at its left edge** (toggle with the *Labels* checkbox; an optional corner *Legend* is also available). A **theme** selector switches between two saved presets:

- **White (publication)** (*default*) — white background, **point + line** markers, a colour-blind-friendly categorical palette (matplotlib *tab10*), a boxed frame and a light grid — ready to drop into a paper.
- **Dark** — the classic on-screen look (dark background, line-only, vivid golden-angle colours).

An x selector plots the axis in **R (px)**, **2 θ (°)** or **Q ($\text{\AA}^{-1}$)** (converted from the run's calibration). The **Grid** checkbox (off by default) toggles the horizontal + vertical grid. The *spacing* box shifts successive curves vertically (0 = overlay). Top-right controls adjust the **line width** (- line +), **symbol size** (- sym + , for the point+line markers — also turns markers on if a line-only theme is active) and **label font size** (- font +).

8. Tab 5 — Corrections & Physics (*work in progress*)

Preview physics corrections on a single frame (auto-loads on browse). Pixel-domain: polarization, solid angle, empty subtraction. Profile-domain: cylindrical absorption (μR), Compton subtraction (composition:fraction). **Compute corrected profile** overlays corrected vs uncorrected and shows the correction factor vs 2 θ .

Also hosts **per-pixel gain training (LearnableGain)**: from a clean reference frame and a drifted frame, learn a spatial gain map $g_i = 1 + \text{scale} \cdot r_i$ by minimising $\text{MSE}(\text{profile}) + \text{unity} \cdot \Sigma(g-1)^2 + \text{smooth} \cdot \text{TV}(g)$; save as NPZ and apply with $\text{corrected} = \text{raw} / \text{gain_map}$.

9. Tab 6 — PDF Analysis (*work in progress*)

Polyatomic **total-scattering** reduction: $I(Q) \rightarrow$ Faber-Ziman structure function $S(Q) \rightarrow$ pair-distribution $G(r)$, powered by the `midas_pdf` backend (real composition-weighted normalization, Compton subtraction, end-to-end σ propagation, and optional differentiable scale/background refinement). This replaces the earlier monoatomic, composition-free $F(Q)=Q \cdot (I/bg-1)$ approximation.

Background — what $I(Q)$, $S(Q)$ and $G(r)$ mean

PDF analysis is three transforms that move from *what the detector measured* to *where the atoms are*. They are exactly the three stacked plots (top $\rightarrow$ bottom).

- **$I(Q)$ — measured scattered intensity.** Intensity vs. momentum transfer $Q = 4\pi \cdot \sin(\theta)/\lambda$ ($\text{\AA}^{-1}$), a wavelength-independent rescaling of the angle 2θ ; high Q = wide angle = fine spatial detail. Obtained by azimuthally integrating the detector rings to a 1-D curve. $I(Q)$ is **not** pure structure — it mixes the interference (coherent) signal with big smooth backgrounds: per-atom self-scattering (form factors $\langle f^2 \rangle$) and inelastic **Compton** scattering.
- **$S(Q)$ — structure function.** $I(Q)$ with the self-scattering and Compton removed and normalized away (the **Faber-Ziman** step), leaving only interference: $S(Q) = [I_{\text{coh}} - (\langle f^2 \rangle - \langle f \rangle^2)] / \langle f \rangle^2$, where $\langle f \rangle$, $\langle f^2 \rangle$ are composition-weighted atomic form factors (hence the **Composition** input). $S(Q)$ oscillates about 1 and $\rightarrow 1$ at high Q where correlations wash out (the dashed guide line). $F(Q) = Q \cdot [S(Q)-1]$ is the same thing re-weighted by Q to emphasize the structurally rich high- Q oscillations.
- **$G(r)$ — reduced pair distribution function.** The real-space answer: a histogram of interatomic distances, via a sine Fourier transform $G(r) = (2/\pi) \cdot \int Q \cdot [S(Q)-1] \cdot \sin(Qr) dQ$. A **peak at r** means many atom pairs are separated by that distance; the **first peak is the nearest-neighbour bond** ($\text{Ni} \approx 2.49 \text{ \AA}$, CeO_2 $\text{Ce-O} \approx 2.34 \text{ \AA}$), later peaks are further coordination shells. Below the first bond $G(r) = -4\pi\rho_0 r$ (nothing can be closer) — a constraint that needs the number density ρ_0 . Because it uses *total* scattering (Bragg and diffuse), the PDF captures local structure even in disordered / nanoscale / amorphous materials, unlike a Bragg-only Rietveld fit.

detector image / $I(Q)$ file

| azimuthal integration

▼

$I(Q)$ raw intensity = interference + form factors + Compton (top plot)

| subtract Compton, subtract Laue $(\langle f^2 \rangle - \langle f \rangle^2)$, divide by $\langle f \rangle^2$ (needs composi

▼

$S(Q)$ pure interference, oscillates about 1 (middle plot)

| Fourier sine transform over $[Q_{\text{min}}, Q_{\text{max}}]$ with a window (Lorch)

▼
G(r) interatomic distances; peaks = coordination shells (bottom plot)

The $\pm 1\sigma$ band on G(r) is the counting uncertainty σ propagated from I(Q) through every step, so real peaks can be told from noise. The **G/g/T/R** family (Output selector) is the same information re-weighted: **G** oscillates about 0 (refinement standard); **g** $\rightarrow$ 1 at large r; **T** is the total correlation; **R** is the radial distribution whose *peak area = coordination number*. g/T/R all need ρ_0 .

I(Q) source (choose one):

- **Integrate detector frame** — a calibrated frame is integrated (hard/polygon binning, Poisson σ) and mapped to Q, exactly as before. Uses the Tab 2 calibration (λ , geometry) and any Tab 1 mask.
- **Load I(Q) file** — a pre-integrated 2- or 3-column Q, I, σ text/CSV (comment/header lines tolerated). The tab **opens on this source by default**, pointed at real data in `test_data/pdf_real/` (`iq_Nickel.csv` ; also CeO₂, IPA, Kapton at λ 0.1839) — just press **Compute G(r)** for the Ni PDF (first peak ~ 2.49 Å). Synthetic known-answer examples are in `test_data/pdf_synth/` (`nickel_iq.csv` , `ceo2_iq.csv`). See each folder's `README.md` for per-sample composition / ρ_0 / Q-range settings.

Calibration. The geometry/wavelength comes from Tab 2 automatically, or use **Load calibration file...** to read a MIDAS `paramtest.txt` , a `.json` , or a pyFAI `.poni` (auto-detected). This is required for *Integrate detector frame* and also sets λ used by *Load I(Q) file* mode.

Sample. Composition (e.g. `Ni` or `C:3,H:8,O:1` ; ions like `Ni2+` allowed), number density ρ_0 (atoms/Å³; needed for refinement and for g/T/R), wavelength λ (auto-filled from calibration), and a **Compton subtraction** toggle.

Normalization. Optional **Refine scale + background** (L-BFGS) — reveals a background polynomial degree (0–3), a low-r cutoff `r_min` (Å), and an iteration count. Refinement requires $\rho_0 > 0$. It fits scale and a smooth b(Q) against two model-free constraints: high-Q $\langle S \rangle \rightarrow 1$ and $G(r) = -4\pi\rho_0 r$ below `r_min` .

Output. Bottom-plot convention family — **G(r)** (reduced PDF), **g(r)** (pair distribution), **T(r)** (total correlation), or **R(r)** (radial distribution, whose peak integral is the coordination number); g/T/R need ρ_0 . Middle plot toggles between **S(Q)** (with the S=1 guide) and the true reduced structure function **F(Q)=Q(S–1)**.

Right panel: I(Q) (+ refined background overlay), S(Q)/F(Q), and the selected G/g/T/R with a shaded $\pm 1\sigma$ band. **Save G(r)** writes a three-column `r, G(r),  $\sigma$` file (diffpy-CMI / PDFgui / RMCProfile compatible); **Save S(Q)** writes `Q, S(Q)` .

Functions behind the tab

The tab is a thin UI over a background worker and the `midas_pdf` backend.

UI — PDFTab (midas_gui/tab_pdf.py)

Method	Role
<code>_build_ui</code>	Builds the I(Q)-source / Sample / Normalization / Output groups and the three stacked plots.
<code>set_calibration(result, source)</code>	Receives a Tab-2 result (or a loaded geometry) → sets λ and enables Compute.
<code>_load_calib_file</code>	Loads a <code>paramstest / .json / .poni</code> → builds a calibration result → <code>set_calibration</code> .
<code>_load_img</code>	Loads a detector frame for <i>Integrate detector frame</i> mode.
<code>_run</code>	Validates inputs, assembles the <code>cfg</code> dict, and starts a <code>PDFWorker</code> .
<code>_on_done</code> / <code>_redraw_mid</code> / <code>_redraw_bottom</code>	Draw I(Q)+bg, the S(Q)↔F(Q) toggle, and the G/g/T/R curve with its $\pm 1\sigma$ band.
<code>_save_gr_file</code> / <code>_save_sq_file</code>	Export <code>r,G,σ</code> and <code>Q,S</code> .

Worker — PDFWorker (midas_gui/workers.py) runs off the GUI thread:

Function	Role
<code>_acquire_iq</code>	Produces <code>(q, I, σ)</code> — either by integrating the frame (image mode) or by reading a file (<code>_load_iq_file</code> , tolerant of comma/space and 2- or 3-columns).
<code>run</code>	Trims to <code>[Qmin,Qmax]</code> , builds <code>Composition</code> , calls the reduction (plain or refine), computes <code>F(Q)</code> and the G/g/T/R family, emits the result dict.

Image-mode integration reuses the shared core `_build_spec`, `build_geom`, `integrate_frame`, `axis_conversions` (same path as the other tabs); geometry-file parsing uses `geometry_fields_from_file` / `result_ns_from_geometry_file` (`midas_gui/helpers.py`).

Reduction — midas_pdf via midas_gui/pdf_backend.py (re-exported symbols):

Function	Does	Maps to plot
<code>Composition(fractions, number_density)</code>	Composition layer: <code>{f}</code> , <code>{f²}</code> form-factor averages, Laue term <code>{f²}-{f}²</code> , and per-atom Compton.	—
<code>faber_ziman_S(I, q, comp, ...)</code>	I(Q) → S(Q) with σ (subtract Compton/Laue, divide by <code>{f}²</code>).	S(Q)
<code>structure_function_F(q, S)</code>	$F(Q) = Q \cdot (S-1)$.	F(Q)

Function	Does	Maps to plot
<code>fourier_sine_transform(q, S, r, window)</code>	$S(Q) \rightarrow G(r)$ with σ (Lorch / none window).	$G(r)$
<code>i_of_q_to_Gr(q, I, comp, r, ...)</code>	End-to-end $I(Q) \rightarrow G(r)$ (composes the two above); the default Compute path.	$S(Q) + G(r)$
<code>refine_normalization(q, I, comp, r, ...)</code>	L-BFGS fit of scale + background polynomial against high-Q $\langle S \rangle \rightarrow 1$ and low-r $G = -4\pi\rho_e r$.	refined $S(Q)/G(r) + bg$
<code>pair_distribution_g /</code> <code>total_correlation_T /</code> <code>radial_distribution_R</code>	Convention family from $G(r) + \rho_e$.	$g/T/R$

Packaging note. `midas_pdf` is currently **vendored** in `midas_gui/_vendor/` and loaded through `midas_gui/pdf_backend.py`, which also installs a small `midas_hkls.absorption` compatibility shim for `midas-hkls < 0.5.0`. Once `midas-pdf` is published and `midas-hkls >= 0.5.0` is available the vendored copy and the shim are removed.

10. Tab 7 — Texture / Pole Figure (*work in progress*)

Per-ring azimuthal analysis for preferred orientation. Controls: calibration source, sample frame, R/η bins, ring index, χ (tilt) / ϕ (rotation). **Compute Pole Figure** integrates to a cake and extracts $I(\eta)$ at the selected ring; the right panel shows the stereographic pole figure and the raw $I(\eta)$.

Save pole figure (.pol) exports POPLA format.

11. Pump Probe (time-resolved / TR-XRD)

Analyses time-resolved (pump-probe) diffraction the way the TRR group does: a folder of raw detector frames is pooled by a filename prefix, the pump-probe **delay** is parsed from each name, every frame is integrated to $I(q)$ with the **MIDAS engine** (the same core as Batch Integrate, driven by a calibration), repeats at each delay are averaged, and a reference (mean of the pre-time-zero / negative delays) is subtracted to give $\Delta I(q, \text{delay})$.

File-naming / pooling. Frames follow `PREFIX-<fshw>fshw<delay>delay<id>.tif`, where `PREFIX` (e.g. `Ex01_Sa01_Sc17` = Experiment / Sample / Scan) is one opaque grouping key. **Scan folder** globs `PREFIX*.tif` in the loaded folder, parses `fshw` and `delay` (seconds) from each name, and reports how many frames and unique delays were found. The delay sign is flipped on load so positive delay = after the pump.

The tab uses the same three-panel layout as Batch Integrate: **data loader** (left), **settings** (middle), **plots** (right).

Left — data loader (the shared loader panel, as on Calibrate / Batch):

- **Data** — the folder of raw frames. Defaults to the bundled TRR test set in `test_data_pump_probe/detimages/` (125 Pilatus2M frames; a large, git-ignored local asset), so the tab is populated on open. An index range / stride can subset the pooled frames.
- **Dark / Bright / Background** — per-frame field corrections (compute each field, then it is applied to every frame before integration).
- **Mask** — a mask file and/or the mask from the Mask Builder tab. For the TRR data the tab pre-loads `invert_mask.tif` (0 = valid pixel, 1 = bad pixel / module gap), which matches MIDAS's convention that non-zero pixels are masked out.

Middle — settings.

- **Calibration source** — *From Tab 2* (the live calibration) or *From file* (`calibration.json` / `paramstest.txt` / `.poni`). Same convention as Batch Integrate. For the shipped TRR test data the field defaults to a ready MIDAS `paramstest` (`Ex01_Sa01_Sc17_midas.txt` , converted from the dataset's pyFAI/Fit2D geometry), so the tab integrates out of the box.
- **Data pooling (TRR)** — the pooling **prefix + Scan folder** (the folder itself comes from the loader on the left).
- **Integration** — kernel, R/ η bins, the plot axis (Q / 2θ / R), and optional Q-uniform binning. Identical engine and options to Batch Integrate.
- **Physics corrections** — polarization / solid-angle.
- **Pump-probe options** — the reference-delay set (defaults to all negative delays; multi-select to override), optional per-pattern normalization over a q-window, the ΔI colour range (auto or fixed $\pm$) and the diverging colormap.

Views (right panel). Every axis shows **real physical values**, not indices — the radial axis is in Q ($\text{\AA}^{-1}$), 2θ ($^\circ$) or R (px) per the plot-axis selector, and delays are in seconds.

1. **ΔI heatmap** — ΔI over the radial axis (a true, continuous Q/ 2θ /R axis) versus delay, diverging colour centred at zero, with a labelled ΔI colour-bar and the reference $I(q)$ lineout beside it (shared radial axis). Delays span several decades and are irregular, so they are laid out as columns labelled with their real delay values.
2. **ΔI vs q** — one curve per delay, coloured along a rainbow by delay. With ≤ 8 delays each curve is named in the legend; with more, a continuous **delay colour-bar** (min $\rightarrow$ max, in seconds) replaces the legend.
3. **Kinetics** — ΔI versus delay for user-defined q-bands (**Add band** over a q range); x-axis as real delay (**linear**), **log** (post- t_0 delays, natural for the decade-wide delay range), or **rank**.
4. **Mean patterns** — the averaged $I(q)$ at each delay plus the dashed reference, with an optional **$\pm 1\sigma$ band** (spread of $I(q)$ across delays) and a **Log Y** toggle to reveal weak

features across the full dynamic range; a signal-level / stability check.

All views use a publication-style white theme with large, clearly-labelled axes, readable legends and a subtle grid. A toolbar above the plots sets the **draw** mode (lines / lines+points / points) and the **line**, **point (sym)** and **font** sizes (the $-/+$ groups), as on the Batch Integrate tab. Pan and zoom are bounded to the data so you cannot lose the plot area.

Integration runs off the GUI thread (`PumpProbeWorker`) with a progress bar and **Abort**. There is no peak fitting — peak position / width / area are read from the plots, matching the reference workflow.

12. Results & Export (*work in progress*)

Session summary + one-click export. Checkboxes select which products (calibration.json, paramtest.txt, mask.tif, integrated profiles, $G(r)$, pole figures, session log) to copy to an output directory. A provenance block (package versions, geometry hash, mask fraction, correction flags) can be copied to the clipboard for a Methods section.

13. Common UI Conventions

- **Browse = load.** Selecting a file/folder (or pressing Enter in a path field) loads it immediately; there are no separate "Load" buttons. HDF5 dataset / frame-index changes reload automatically.
- **No accidental scroll changes.** Spin boxes and drop-downs ignore the mouse wheel — values change only by clicking/typing; the wheel scrolls the panel instead.
- **Readable right-click menus.** pyqtgraph plot context menus use the dark theme.
- **Dark theme, orange accent** (Dioptas-inspired); off-white text, light input fields.
- **Sample-to-detector distance is entered/shown in mm** (Data Viewer, Calibrate seed, Preferences ▶ Geometry). Internally — all calculations — and in written calibration files (`paramtest.txt` , `calibration.json`) it is always in **microns**; the $\text{mm} \leftrightarrow \mu\text{m}$ conversion happens only at the display boundary. The config key remains `lzd_um` (μm).

14. Packaging, Deployment & Diagnostics

Packaging. `midas-gui` is a MIDAS-style package: `pyproject.toml` (BSD-3-Clause), a `tests/` smoke suite, and `release.sh` for cutting versioned releases (see `RELEASING.md`). Once published it can join the `midas-suite` meta-package as an optional `gui` extra (`pip install midas-suite[gui]`).

Deployment on another system. Clone the repo, create the conda environment (`environment.yml`), and run `midas-gui` . The MIDAS analysis backends are installed via the `pip:` section of `environment.yml` (editable from a local MIDAS checkout, from a git subdirectory, or from a private index). `midas_gui/_paths.py` is an inert stub in an installed package (no `sys.path` manipulation).

Crash diagnostics. On startup the app installs a logging exception hook and `faulthandler` ; any uncaught Python exception or native fault is written to `~/midas_gui_error.log` and shown in a dialog. Installing the hook also prevents PyQt5 from hard-aborting on an exception inside a slot. Each tab is built in isolation — a tab that fails becomes an error placeholder instead of taking the window down. If the app ever "pops up and dies" (typically on Windows), send that log file.

15. Configuration & Defaults

Every default in the GUI — detector geometry, data/output paths, ring-simulation **materials**, **calibrants**, the **pixel-preset** and **K-edge** menus, the calibration / integration **algorithms**, and **which tabs are visible** — can be set **without editing code**, via a single per-user JSON config that overrides the shipped built-in defaults.

Where it lives

One per-user file, auto-located per OS: `~/.config/midas_gui/config.json` (Linux), `~/Library/Application Support/midas_gui/config.json` (macOS), `%APPDATA%\midas_gui\config.json` (Windows). If it's absent, the shipped built-in defaults are used. Sharing between machines/users is done by exporting/importing a JSON file (below) — copy it wherever you like.

Profiles — Settings ▶ Preferences... ▶ Profile row

The top of the Preferences dialog has a **Profile** row — `Profile:` `[combo ▼]` `[New...]` `[Duplicate...]` `[Rename...]` `[Delete]` — for keeping several independent sets of defaults (e.g. one per beamline) and switching between them:

- Each profile is its own config file under `<config_dir>/profiles/<name>.json` ; the active one is remembered in `<config_dir>/profile_meta.json` . Existing single-config installs are migrated transparently into a profile named **Default** the first time this runs — no data is lost.
- **New...** seeds a blank profile from the shipped built-in defaults.
- **Duplicate...** seeds a new profile from whatever is currently shown in the dialog (including unsaved edits), then switches to it.
- **Rename... / Delete** — Delete refuses to remove the last remaining profile; deleting the active profile switches to another one first.

- Switching profiles (via the combo) prompts to confirm if the dialog has unsaved edits, then reloads the dialog's fields and the live `DEFAULT_*` globals from the new profile — most values apply immediately; a few (e.g. viewer step sizes set at widget-construction time) need a restart, and you're offered one.
- The "Local config: ..." label under the row becomes "**Profile '<name>': <path>**" so you always know which file you're editing.

Editing — Settings ▶ Preferences...

The **Settings** menu opens **Preferences...**, a dialog whose tabs are **pre-filled with the full shipped defaults** so you edit from a complete starting point:

- **Geometry** — λ , pixel size, Lsd, beam centre.
- **Data Viewer** — the up/down-arrow step size for each field in the Data Viewer's Ring simulation card (λ , max 2θ , Lsd, pixel size, BC_y/BC_z, ty/tz). Shipped defaults: 0.01 Å, 1°, 1 mm, 0.1 μm, 1 px, 0.1° respectively.
- **Paths** — default data / calibration / output files & folders.
- **Materials / Calibrants** — add / remove / modify (name + lattice + SG).
- **Devices** — the detector devices offered in the Data Viewer's **Live Data** PV dropdown (name, prefix, PVA suffix). The live PV is built as `prefix + PVA suffix`. Ships pre-filled with the 20-ID-D detectors (`oryx20idd` , `s20idPil` , `pg4` , `gh1s` , `s20varex1`), all with PVA suffix `Pva1:Image` ; add / remove / edit rows for your own beamline's devices.
- **Menus** — the pixel-size presets and K-edge foils.
- **Algorithms** — default calibration pipeline, integration kernel, output format, error model, colormap/theme.
- **Tabs** — which tabs are visible. Data Viewer, Mask Builder, Calibrate and Batch Integrate are always shown (their boxes are locked on); tick/untick the rest. **Tab visibility applies immediately** on save (no restart), unlike the other settings.
- **Display** — the **interface scale**, a whole-application zoom (layout *and* fonts) for HiDPI / 4K monitors where the default looks too small. Pick a value or a preset (100 % $\approx$ 1080p, 150 % $\approx$ 1440p, 200 % $\approx$ 4K); it applies after a restart (you're offered one on save). Also reachable from **Settings ▶ Interface scaling....** It is implemented with Qt's `QT_SCALE_FACTOR` , so the whole layout scales uniformly.

Buttons: **Save as my defaults** (write your config), **Save current GUI state** (capture the Data Viewer's live λ /pixel/Lsd/BC into the fields), **Save config to JSON... / Load config (JSON)...** (export/import a config to share), and **Reset to shipped defaults** (delete your config). Also on the menu: **Open config folder** and **Reload config**. Saving writes your per-user file; **changes take effect on the next launch**.

File format

```
{
  "geometry": { "wavelength_A": 0.39, "pixel_um": 75.0, "lsd_um": 121000.0,
    "bc_y": 10.0, "bc_z": 10.0,
    "pixel_presets": [ ["Eiger", 75.0], ["Pilatus", 172.0] ],
    "k_edge_foils": [ ["Au", 80.725], ["Pb", 88.005] ] },
  "viewer_steps": { "wavelength": 0.01, "two_theta": 1.0, "lsd_mm": 1.0,
    "pixel": 0.1, "bc": 1.0, "tilt": 0.1 },
  "materials": { "Ni (FCC)":
    { "a": 3.5238, "b": 3.5238, "c": 3.5238, "alpha": 90, "beta": 90, "gamma": 90, "sg": 225 }
  },
  "calibrants": { "CeO2":
    { "a": 5.4116, "b": 5.4116, "c": 5.4116, "alpha": 90, "beta": 90, "gamma": 90, "sg": 225 }
  },
  "devices": [ { "name": "s20varex1", "prefix": "20IDFF:", "pva_suffix":
    "Pva1:Image" } ],
  "paths": { "nickel_h5": "/data/mygroup/sample.h5", "calib_file":
    "/data/mygroup/calibration.json" },
  "ui": { "calibration_pipeline": "one_shot", "integration_kernel": "subpixel2",
    "output_format": "csv", "azimuthal_method": "poisson", "plot_theme":
    "hot",
    "visible_tabs": ["Calib. Refinement", "Pump Probe"],
    "ui_scale": 1.0 }
}
```

- `ui.visible_tabs` lists the **optional** tabs to show (the four always-on tabs are implicit). The shipped default is `["Calib. Refinement", "Pump Probe"]` — Corrections, PDF Analysis, Texture and Results & Export are hidden until you add them. Edit it from **Preferences ▶ Tabs**.
- `ui.ui_scale` is the whole-interface zoom (0.5–4.0) applied at startup via `QT_SCALE_FACTOR`; ~1.5 for 1440p, ~2.0 for 4K. Edit it from **Preferences ▶ Display** or **Settings ▶ Interface scaling...** (takes effect on restart).
- When a **list section is present it fully replaces** that built-in list — so `materials`, `calibrants`, `devices`, `pixel_presets` and `k_edge_foils` in the file are your complete lists. (The Preferences dialog pre-fills them with the shipped entries, so you always start from the full set; use **Reset to shipped defaults** to get them back.) `sg` is the space-group number.
- Geometry scalars, `paths` and `ui` are simple overrides; any section or key may be omitted. A malformed config is ignored (built-ins are used) rather than blocking startup.
- A minimal template lives at `documentation/config.example.json`; the step-by-step guide is `documentation/config_gui.md`.

16. File ▶ Save/Load GUI State

Beyond a saved **profile** of defaults (§15), the **File** menu can capture the **live, in-progress state of every tab** — every field you've typed or picked, across all 10 tabs — to one JSON file, so a session can be closed and resumed later exactly where it left off.

- **File ▶ Save GUI State...** (`Ctrl+S`) — the first time in a session, prompts for a destination (default `midas_session.json`); every subsequent `Ctrl+S` **silently overwrites that same file** (no dialog) as long as the session hasn't loaded/saved a different file since. **File ▶ Save GUI State As...** (`Ctrl+Shift+S`) always prompts for a destination, and makes *that* file the target for future plain `Ctrl+S` . Either way, a completion dialog lists which tabs (if any) failed to save.
- **File ▶ Load GUI State...** (`Ctrl+O`) — asks for confirmation (loading overwrites every tab's current values), then pick a previously saved file. A completion dialog reports any tabs present in the file that no longer exist, or that failed to restore. The tab that was active at save time is re-selected.

What is restored automatically

Every field (text boxes, spin boxes, checkboxes, combo/dropdown selections) in every tab is restored as typed. In addition, **path-backed data is reloaded from disk**, the same way it loads when you type/browse to a path by hand — images, masks-by-path, dark/bright/background frames, and HDF5 datasets all re-read their file automatically after a state load (guarded so a moved/deleted file is skipped quietly rather than popping a warning).

What is *not* re-run

Loading a state **does not re-run any long-running pipeline** — Calibrate's Fit, Batch Integrate, the PDF transform, and Calibration Refinement all keep their inputs restored but require **one manual click of the tab's own Run/Fit button** to reproduce their result. This keeps a load fast and avoids silently kicking off a multi-minute job in the background.

Sidecar files for in-progress derived data

Two tabs can hold computed data that hasn't been exported to a file of its own yet — a drawn/computed **mask** (Mask Builder) and a just-fit **calibration result** (Calibrate). Saving a GUI state writes small sidecar files next to it so this in-progress work isn't silently lost:

- `<state file stem>_mask.tif` — the current in-memory mask, if any; reloaded automatically on the next Load.
- `<state file stem>_calibration.json` — the last fit result's flattened fields, kept for the record and to reseed the manual/seed geometry fields. This does **not** reconstruct the in-memory fitted result object (it isn't a plain, re-loadable structure) — re-run **Fit** after loading to reproduce it.

For bugs or questions, see the MIDAS GUI repository.